Agent Feedback & Analytics — Stakeholder Guide

How users rate the AI assistant, and how to read every number on the Assistant Analytics dashboard. *(Written for business readers — no technical background needed.)*

1 · How Users Give Feedback

1. **After a help answer** — a row of 5 stars appears under every reply. One click submits; a comment box then opens if the user wants to say more.
2. **After creating a service request** — the user is asked once, next to the **View Service Request** card. The whole flow is rated, not each question the agent asked along the way.
3. **Comments are optional.** Users rating 3 stars or below are gently prompted with *"What was missing?"*
4. Ratings are always tied to the signed-in user, so we know who said what and can follow up.

2 · Where to See It

Open **Dashboard** → **Assistant Analytics**. The page has two tabs:

- **Overview** — how the assistant is serving tenants. For business and operations.
- **Agent health** — how the assistant is performing technically. For the engineering team.

3 · Overview Tab — What Each Number Means

Performance cards

Card	What it tells you
Response Accuracy	Of the questions users asked, the share the assistant actually answered.
Questions Asked	How many real questions users typed. Background system checks are excluded.
Peak Concurrent Users	The busiest hour's number of active conversations.
Avg Response Time	How long the assistant takes to reply, in seconds.
Escalation Rate	The share of questions the assistant could not answer — the ones a person has to pick up.
Total Sessions	How many separate conversations took place.

Summary strip

Figure	What it tells you
Questions Not Answered	The percentage the assistant failed to answer. The mirror of accuracy.
User Satisfaction (CSAT)	The average star rating, shown as a percentage of 5 stars.
Avg Session Length	How long a typical conversation lasts, in minutes.
Fallback Rate	How often the assistant replied "could you rephrase that?" instead of answering.
Questions per Session	How many questions a user asks in one sitting.

Cards below

Card	What it tells you
Sessions by Hour	When during the day people use the assistant. Useful for staffing and release timing.

Card	What it tells you
User Ratings	The average score, how many people rated, and the spread from 1 to 5 stars. The spread matters: a 3.5 built from all 3s is a very different problem from one built from 5s and 1s.
Top Users	Who uses the assistant most, and the rating each of them gives. A heavy user with a low score is the most important row on the page.
Top Questions	What people ask most. A red **"N unanswered"** badge marks questions the assistant repeatedly fails — the first place to improve the help content.
Recent Conversations	The latest conversations, with the opening question, any comment left, the agent that handled it, and the rating.

4 · Reading an Actual Conversation

1. Click any **session ID** in the conversations list.
2. You see the **full exchange** — every question, the reply it received, and the rating and comment on each one.
3. Ratings sit next to the answer that earned them, so a one-star with *"no idea"* can be traced to exactly what the assistant said.

5 · Agent Health Tab — What Each Number Means

Section	What it tells you
Slowest Step	The part of answering that takes longest. Where to look if replies feel slow.
Least Reliable API	The platform service that fails most often when the assistant calls it.
Biggest Token Consumer	Which part of the assistant uses the most AI processing — the main driver of running cost.
Where the Time Goes	Every step ranked slowest first, so delays are visible at a glance.
Platform API Reliability	Success rate for each platform service the assistant depends on. When a request stalls, the cause is usually here rather than the AI.
Token Spend by Prompt	Which parts of the assistant consume the most processing, and how much.
Where Service Requests Stall	How many conversations reach each stage of creating a request — showing exactly where people drop off.

6 · How Response Accuracy Is Worked Out

Of the questions users asked, the share the assistant answered.

> **Example** — 100 conversation turns. 30 are background system checks, leaving **70 real questions**. The assistant answered 63 of them.
>
> **Accuracy = 63 ÷ 70 = 90%**

Two supporting figures always travel with it:

- **Sample size** — how many questions the percentage covers.
- **Coverage** — that sample as a share of all questions asked.

This matters: *100% accuracy* means very little if it is measured over two questions out of eighty. Always read the percentage together with its coverage.

The assistant now tells us how it is doing in users' own words — and every score links back to the answer that earned it.